

Darsh Bhatt

(647) 219-6049 | darshb.08@gmail.com | www.linkedin.com/in/darshb08

PROFILE

AI Engineer building LLM applications, agentic workflows, RAG systems, and Python-based data platforms across enterprise and research environments. Experienced across end-to-end AI development spanning data ingestion, retrieval, context engineering, model and tool integration, evaluation, APIs, and containerized deployment using Python, LangGraph, Docker, Kubernetes, and cloud infrastructure. Strong at translating ambiguous business and technical problems into practical, reliable AI solutions.

CORE COMPETENCIES

Languages: Python, Java, R, SQL, Shell Script

Applied AI / LLMs: LangGraph, RAG, Agentic Workflows, MCP, Vector Databases, Context Engineering, Prompt Engineering, Function Calling, Structured Outputs, LLM Integration

AI Engineering: Retrieval Evaluation, LLM Evaluation, Token Optimization, API Integration, Workflow Orchestration

ML / NLP: BERT, PyTorch, TensorFlow, spaCy, Scikit-learn

Cloud / Infrastructure: AWS (S3, EC2), Docker, Kubernetes, High-Performance Computing (HPC)

Data / Systems: Pandas, NumPy, Databricks, REST APIs, ETL Pipelines, NoSQL, Oracle Autonomous DB

Frontend / Visualization: React, Power BI

Version Control / Tools: Git, GitLab, RStudio, Jupyter, Azure DevOps

EXPERIENCE

Charger Logistics Inc. (AI Engineer, Toronto, Canada)

Aug 2026 - Present

- Developing LLM-powered workflows using LangGraph, RAG, and agentic orchestration for enterprise operational use cases.
- Building context and retrieval workflows that integrate LLMs with internal data sources, APIs, and tools using structured outputs and function calling.
- Evaluating retrieval, prompt, and context strategies to improve grounding, response quality, token usage, and reliability across LLM workflows.
- Working with MCP-based tool integrations to enable LLM agents to securely interact with enterprise systems, services, and external capabilities.
- Supporting Python-based AI services and containerized deployment workflows using Docker and Kubernetes.

University of Windsor (Research Assistant 2, Windsor, Canada)

May 2024 - Aug 2026

- Designed containerized analysis workflows using Docker, Kubernetes, and HPC resources to support compute-intensive bioinformatics processing across multi-node environments.
- Built a RAG-based lab assistant using Claude API and vector search over internal documentation to help researchers access technical context through natural-language prompts.
- Improved retrieval and prompt orchestration workflows to reduce manual support by 80% and email-based knowledge transfer, enabling faster self-serve access to internal technical context across the lab.
- Built and optimized Python, R, and SQL pipelines to process and analyze 1M+ phospho-proteomic records, reducing end-to-end runtime by 50% and improving reproducibility across research workflows.
- Applied machine learning and statistical methods, including clustering, dimensionality reduction, regression, and hypothesis-driven analysis, to high-dimensional biological datasets.

ZenQuip (Applied AI Developer, Part-time, Remote, Canada) | zenquip.com

Jan 2026 - Aug 2026

- Built LLM-based workflows for an in-development mental health platform to support conversational interactions and guided user assistance.
- Developed RAG pipelines using LangGraph, vector retrieval, and internal content to improve response grounding and surface relevant support resources.
- Worked with multiple LLM APIs, including OpenAI, Anthropic, and Groq, to test models, refine prompt behavior, and reduce token usage and cost per request by 40% through adaptive RAG workflow tuning.
- Implemented conversation flow logic based on user input to make interactions more context-aware and better aligned with support needs.

EDUCATION

Master of Applied Computing | University of Windsor, Windsor, ON

May 2023 - Aug 2024

B.Tech in Computer Science and Engineering | Gujarat Technological University, INDIA

Aug 2018 - Mar 2022